

User Guide

Contents

User Guide	2
1. Getting Started	2
1.1 Standalone mode	2
1.2 Platform mode	2
2. Managing Workspaces	3
2.1 Creating a new workspace	4
2.2 Renaming a workspace	4
2.3 Exporting a workspace	5
2.4 Deleting a workspace	5
3. The Editor at a Glance	6
4. Importing Source Text	7
5. Working with Segments	7
5.1 Automatic segmentation	7
5.2 Adjusting segment boundaries	7
5.3 Splitting a segment	8
5.4 Translation gaps and the action guard	8
6. Named Entity Recognition	9
6.1 The segment toolbar	9
6.2 Running NER	9
6.3 Adding an entity manually	10
6.4 Editing or deleting an entity	10
6.5 Resolving conflicts	11
7. Translations	12
7.1 Adding a target language	12
7.2 Switching between layers	13
7.3 Post-editing a translation	13
7.4 Re-translating over edited content	13
7.5 Clearing a translation	14
8. Thesaurus Tagging	14
8.1 Searching the thesaurus	14
8.2 Tags applied to a segment	15
8.3 Adding a free-text tag	16
8.4 Removing a tag	16
8.5 Automatic semantic tagging	16
9. Exporting Your Work	17
10. Administrator: Configuring NLP Endpoints	17

User Guide

The Memorise Curation Tool is a configurable web application for curating textual archives. This guide walks through the curator workflow from first login to export, plus the administrative configuration panel. The platform supports two deployment modes. *Standalone* runs entirely in the browser without a backend. *Platform* adds a backend gateway and multi-user accounts. The two modes differ only at the authentication and persistence boundaries; the curation workflow itself is identical in both.

1. Getting Started

1.1 Standalone mode

In standalone mode the application runs entirely in the browser and persists workspaces in the browser’s `localStorage`. The login screen offers only a sign-in form, because there is no shared backend that would distinguish accounts beyond the device.



Figure 1: Login screen in standalone mode

A first-time user signs in with a chosen local username; the application creates the local profile on the fly and treats subsequent logins with the same username as the same user on that device.

1.2 Platform mode

In platform mode the application connects to the backend gateway. The login screen exposes both a sign-in and a registration tab, and accounts are persisted server-side so the same user can log in from any machine.



Figure 2: Login screen in platform mode

To create a new account, switch to the registration tab and supply a username, an optional email, and a password.



Figure 3: Registration tab

The username must be between 2 and 50 characters, the password between 4 and 128 characters, and the email, when provided, must be well-formed. The system rejects duplicate usernames with an inline message. A successful registration signs the new user in immediately.

2. Managing Workspaces

After login the workspace list is the curator's home screen. Three example workspaces are pre-seeded on a fresh install so a new curator can explore the editor without first having to import their own text.


Name	ID	Last Updated	Action
Workspace A 	7562f905-d3bd-4ecb-b4f7-870da5a6ca8b	5/3/2026, 8:58:58 PM	<button>OPEN</button>  
Workspace B 	3c6a3bb3-c3db-48c7-8a96-acff9a6ace92	5/3/2026, 8:58:58 PM	<button>OPEN</button>  
Workspace C 	d2d5ec7e-1824-45e5-bc98-460b03347a63	5/3/2026, 8:58:58 PM	<button>OPEN</button>  

Figure 4: Workspace list with pre-seeded workspaces

Each row exposes inline actions: a rename pencil next to the name, an **Open** button, an export icon, and a delete icon.

2.1 Creating a new workspace

The new-workspace control opens a small dialog asking for a name and offering an optional file-import field. The name must be non-empty after trimming and unique among the curator's workspaces. Attaching a `.txt` file (up to one megabyte) seeds the workspace's source text directly; leaving the field empty opens the editor with an empty placeholder into which the curator can paste text afterwards.



The dialog box is titled "New Workspace". It contains a text input field labeled "Workspace name" with the text "New Workspace #2" entered. Below the input field is a file upload button labeled "Upload text file (optional)" with a document icon. At the bottom right are two buttons: "Cancel" and "Create".

Figure 5: Create workspace dialog with the optional file-upload field

After confirmation the system creates the workspace, persists it, and opens the editor on it.

2.2 Renaming a workspace

The rename pencil turns the workspace name into an inline editable field. While typing, the system checks the candidate name against the curator's other workspaces and surfaces a duplicate warning

if one already uses the same name.



Figure 6: Inline rename with duplicate warning

2.3 Exporting a workspace

The export action opens a dialog where the curator chooses between JSON and PDF. Both formats are generated in the browser, so export remains available in standalone mode.




Figure 7: Export workspace dialog

A more detailed account of what each format contains appears in Section 9.

2.4 Deleting a workspace

The delete action requires confirmation, naming the workspace to prevent accidents. Deletion is permanent: in platform mode it removes the workspace from the server, and in standalone mode it removes it from the browser's storage.




Figure 8: Delete confirmation

3. The Editor at a Glance

The editor view is where the curation work happens. It surfaces every concern the curator needs to operate on a single document.

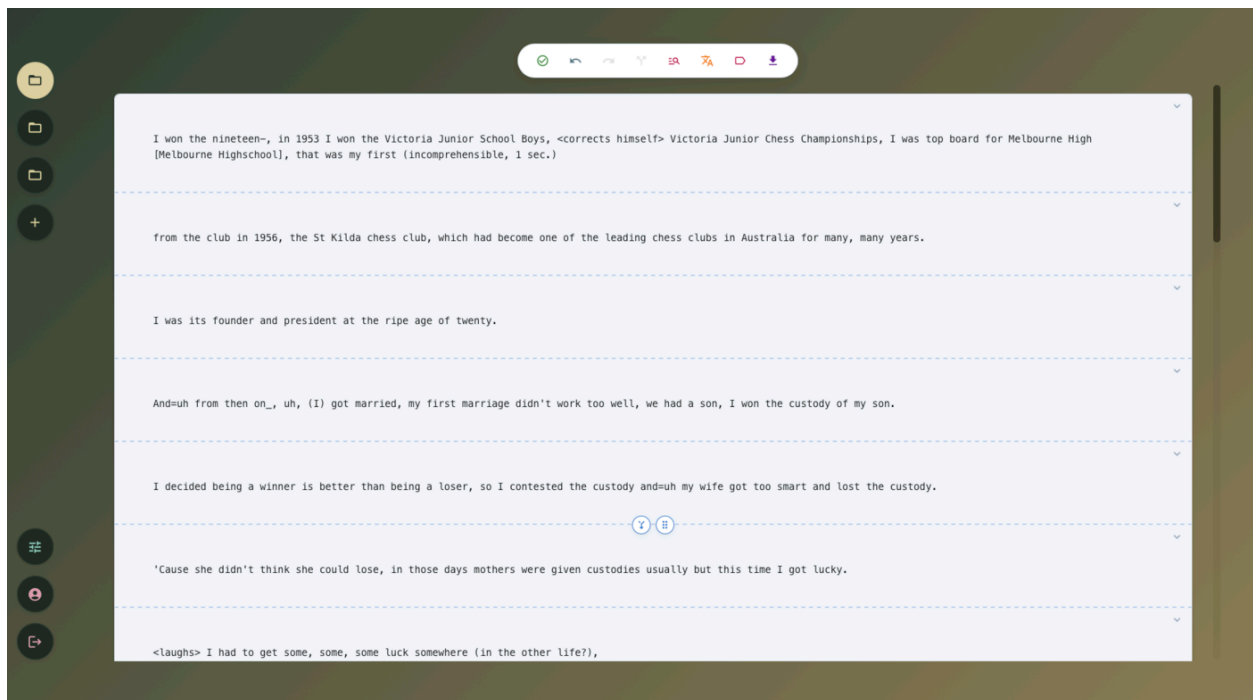


Figure 9: The editor view

The view consists of four regions:

1. **Sidebar (left).** Navigation between workspaces, the new-workspace control, and the logout

- button. Users with the admin role also access the configuration panel from here.
2. **Segmented document editor (centre).** The body of the text, divided into segments after segmentation has been requested, with inline named entity highlights and per-segment toolbars.
 3. **Global action toolbar (top of the editor).** Triggers document-level operations: save, undo and redo, automatic segmentation, named entity recognition, machine translation, semantic tagging, and export.
 4. **Tag panel (right, opened on demand).** Thesaurus and custom tags applied at workspace and segment scope.

4. Importing Source Text

There is no separate import dialog inside an active workspace; source text is supplied at workspace creation. The curator may either attach a plain-text file in the new-workspace dialog (Section 2.1) or confirm the dialog without a file and paste the text into the empty editor placeholder.




Figure 10: Empty editor region for source text entry

Once text is present, all subsequent NLP actions (segmentation, NER, translation, tagging) become applicable to the document.

5. Working with Segments

Segmentation divides the source text into logical units. Annotations are bound to the document by character offsets rather than to individual segments, so segment boundaries can be adjusted later without invalidating named entity or translation alignment.

5.1 Automatic segmentation

The Segment action on the global toolbar dispatches the source text to the configured segmentation service. The response defines the initial segment boundaries, which the editor renders as dashed markers between segments (visible in the editor overview in Section 3).

5.2 Adjusting segment boundaries

Two controls appear between adjacent segments: a **merge** icon on the left and a **drag handle** on the right.




Figure 11: Merge icon (left) and drag handle (right) between two segments

The merge icon collapses two adjacent segments into one. The drag handle, when dragged, shifts the boundary forward or backward to expand or contract the source segment relative to its neighbour.

5.3 Splitting a segment

To split a segment in two, the curator selects a delimiter character (! ? . - , :) inside the segment's body. A small floating menu appears next to the selection offering a *Split segment here* action; confirming it splits the segment at that character position and generates a fresh identifier for the new segment.

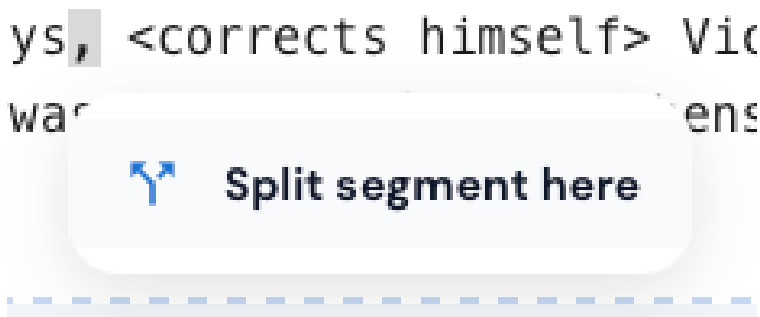


Figure 12: Floating split menu after selecting a delimiter character

The split action is only offered while viewing the original text layer (not on a translation), so that segment-level translation entries cannot fall out of alignment with the original.

After every structural change the system recalculates the character offsets of all downstream annotations so that highlights remain aligned with the words they annotate.

5.4 Translation gaps and the action guard

When a structural edit would leave one language with translations and another without (for instance, shifting a boundary across a segment whose English translation is filled in but whose Belarusian translation is not), the action guard interrupts the operation and asks the curator how to resolve the inconsistency.




Figure 13: Action guard dialog when shifting across an untranslated segment

The dialog enumerates each affected segment and offers to translate the missing one in place before continuing the shift, so the curator never silently loses content.

6. Named Entity Recognition

6.1 The segment toolbar

Each segment carries a header strip that exposes the per-segment controls: the language tab on the left, and the per-segment NER and Tag buttons on the right. The same header also displays inline entity highlights inside the segment body.



Figure 14: Segment with highlighted entities and per-segment header

The two coloured buttons on the right are, in order, the per-segment NER trigger and the per-segment Tag trigger.



Figure 15: Per-segment NER and Tag buttons

6.2 Running NER

NER analysis can be requested per segment from the buttons above, or globally from the document-level toolbar. The system dispatches the relevant text to the configured NER service, and renders returned spans as inline highlights, color-coded by entity type. For global NER the system iterates over every segment and presents a partial-success notice if some segments succeeded and others failed.

6.3 Adding an entity manually

Selecting a range of text in the editor reveals a floating entity-type picker. The picker lists the entity categories supported by the active layer, each marked with the bullet colour used for inline highlights. Choosing a category creates a user-curated annotation at the selected coordinates, distinct from API-generated annotations.




Figure 16: Entity-type picker after selecting a fresh range of text

6.4 Editing or deleting an entity

Clicking an existing highlight opens the same floating menu in edit mode. The top of the menu exposes a text field that lets the curator correct the span's surface form (for instance, extending *Auschwitz* to *Auschwitz-Birkenau*); the category list lets the curator change the entity type; and a **Delete** action at the bottom of the menu removes the span entirely.




Figure 17: Edit menu after clicking an existing entity highlight

Deleting a span records its **start:end:entity** triple in a deleted-set, so re-running NER will not resurrect the same span at the same position and entity type.

6.5 Resolving conflicts

When re-running NER produces a span that overlaps an existing user-curated annotation, or an existing API annotation with a different entity type, the system pauses the operation and asks the curator which version to keep.

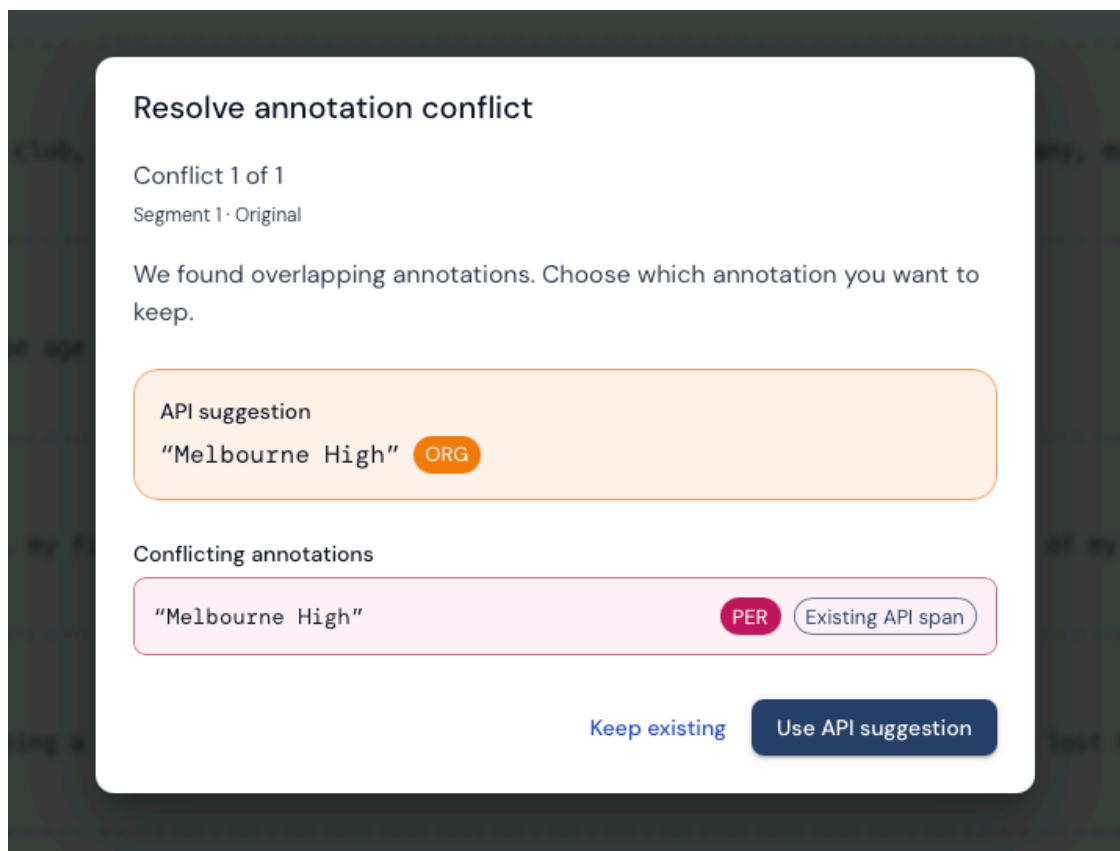


Figure 18: Conflict resolution dialog

API-vs-API overlaps where both annotations agree on the entity type are resolved silently in favour of the newer span. All other overlaps require an explicit decision.

7. Translations

7.1 Adding a target language

The language tab on the left of the segment header exposes the active layer (initially **Original Text**). The plus button beside it opens a language picker for adding a new target language.



Figure 19: Language tab and add-language control

The picker is a searchable list of supported target languages, populated from the configured translate service.




Figure 20: Adding a new target language

7.2 Switching between layers

Once a translation exists, the language tab becomes a dropdown that switches the segment view between the original text and any of its translations.

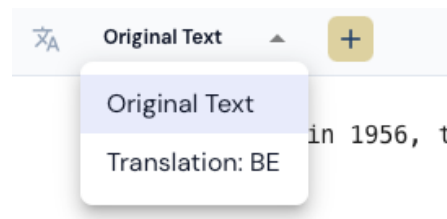


Figure 21: Switching between Original and existing translations

7.3 Post-editing a translation

The translated text region is editable. Saving an edit marks the segment translation as user-edited; the indicator next to the segment header shows that the curator's prose, not the API response, is currently in effect.

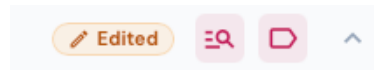


Figure 22: Edited indicator on a translated segment

7.4 Re-translating over edited content

The per-segment re-translate action overwrites the existing segment translation with a fresh API response. If the segment carries the *Edited* badge, the system first presents a confirmation dialog

naming the language and warning that the manual edit will be overwritten. Cancelling the dialog leaves the existing translation untouched; confirming runs the API call, replaces the segment translation, and clears the edit flag.

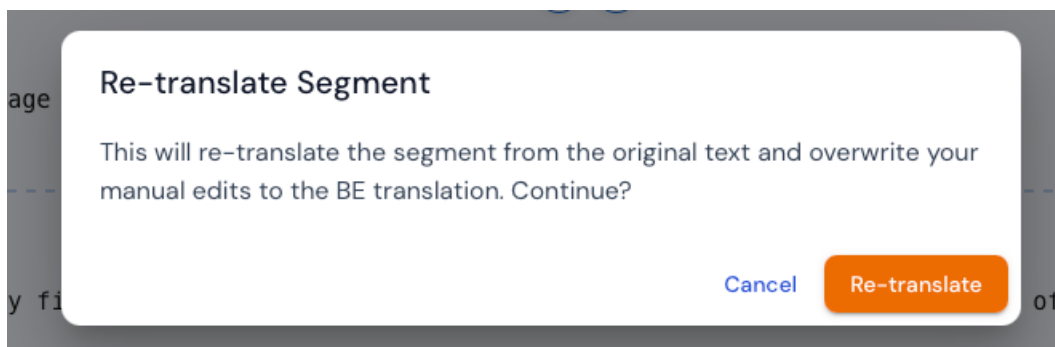


Figure 23: Re-translation warning for an edited segment

The document-level *Translate* action on the global toolbar behaves differently. It only fills in segments that do not yet carry a translation in the target language, leaving any existing translation, including edited ones, untouched. The warning dialog therefore does not apply to that action.

7.5 Clearing a translation

To remove a translation entry for a single segment without affecting other languages, the curator opens the segment header, switches to the target language, and confirms the clear action. The system asks for confirmation before discarding the per-segment entry.

8. Thesaurus Tagging

8.1 Searching the thesaurus

The Tag button opens the tag panel. Its scope tracks the editor selection: with a segment active it shows **Segment Tags** for that segment; with no segment active it shows **Document Tags** for the workspace. The search field at the top of the panel queries the historical thesaurus; matched entries appear in a dropdown with their hierarchical path, and **alias** markers indicate alternative names that resolve to the same canonical entry.




Figure 24: Tag panel with thesaurus search active

The thesaurus search runs in a Web Worker, so the editor remains responsive even on a large thesaurus.

8.2 Tags applied to a segment

Selected tags are grouped by their parent category in the panel, with a count next to each group. Custom (free-text) tags are visually distinct from thesaurus entries.




Figure 25: Segment tags panel populated with selected entries

8.3 Adding a free-text tag

Toggling the **Free Text** switch at the top of the panel changes the input from thesaurus search to a literal-tag input, so the curator can attach an ad-hoc label that is not present in the thesaurus.



Figure 26: Free-text tag mode toggle

8.4 Removing a tag

Hovering over a tag chip in the panel reveals a small **x** button on its right edge. Clicking it removes the tag from the active scope (segment or document) without affecting other segments or other tags. There is no confirmation dialog; the action is reversible through the global undo control (Section 11).

8.5 Automatic semantic tagging

The semantic tagging action (labelled SEM-TAG in the global toolbar) requests a topical classification of a segment or the entire workspace against the configured thesaurus. Suggested tags are

added to the tag panel automatically, scoped either to the relevant segment or to the workspace.

9. Exporting Your Work

The export action is available both from the workspace list (Section 2.3) and from the global toolbar inside the editor. The toolbar exposes the export controls as a small icon group next to the save action.




Figure 27: Export controls on the global toolbar

The format chooser (Section 2.3) offers JSON and PDF. The JSON format mirrors the internal workspace structure, including the deleted-API-key set and the per-segment edit flags, so that the document can be re-imported into another instance without losing state. The PDF format renders the curated content as a human-readable document for archival and sharing. Both formats also contain the per-session curator action counts collected during curation.

10. Administrator: Configuring NLP Endpoints

Users with the administrator role have access to a Services panel from the sidebar. The panel lists every NLP service type with its current endpoint URL, the selected adapter, the latest health-probe result, and the round-trip latency of the last probe.

SERVICES					
API Endpoints		Server (http://localhost:3001)		EDIT REFRESH	
Endpoint	URL	Adapter	Status	Latency	Last Checked
Named Entity Recognition	https://ner-api.dev.memorise.sdu.dk/recogni...	sdu-ner	Up	256ms	10:05:14 PM
Text Segmentation	https://textseg-api.dev.memorise.sdu.dk/seg...	sdu-segment	Up	183ms	10:05:14 PM
Semantic Classification	https://semtag-api.dev.memorise.sdu.dk/clas...	sdu-classify	Up	258ms	10:05:14 PM
Machine Translation	https://mt-api.dev.memorise.sdu.dk/translate	sdu-translate	Up	362ms	10:05:14 PM
Supported Languages	https://mt-api.dev.memorise.sdu.dk/support...	default	Up	361ms	10:05:14 PM

Figure 28: Admin Services panel with endpoint table

The **Refresh** control above the table re-runs every health probe without modifying any configuration, which is useful when investigating intermittent upstream failures.

To change a URL or to switch the adapter for a given service, the administrator clicks **Edit**. The URL becomes editable and the adapter cell becomes a dropdown of adapters registered for that service type. The dropdown lists, for example, the SDU service and the Mock Medical (Comprehend-style) adapter used for the medical-domain reproduction described in the thesis.




Figure 29: Endpoint table in edit mode with the adapter dropdown open

Saving a change writes the updated configuration through the configuration service. The administrator then triggers the Refresh control to re-probe the affected endpoint and verify reachability; the table updates with the new status.

11. Auto-save and Recovery

The auto-save mechanism persists every change without explicit user action. Typing bursts are debounced with a three-second idle window; discrete actions such as confirming a tag or accepting a span use a one-second window; and a thirty-second hard cap guarantees that no session ever drifts further than half a minute from the persisted state. Navigation, logout, and workspace switching trigger an immediate forced save before leaving the page.

The global toolbar additionally exposes **Undo** and **Redo** controls. They roll back individual mutations recorded by the session store, such as span creation, tag attachment, segment merge, or translation edit, and replay them on demand.

If the curator attempts to navigate away or close a workspace with unsaved changes still in flight, the application interposes an unsaved-changes dialog with **Save**, **Discard**, and **Cancel** options, so that no work is lost to an accidental click.

When an external NLP service is unreachable the corresponding action surfaces an error notice rather than failing silently. The curator's existing data is never modified by a failed request, so retrying once the service is back is always safe.